

DynamoDB

定价：
<https://aws.amazon.com/cn/dynamodb/pricing/>

定价概览

DynamoDB 与传统的 NoSQL 部署不同，后者要求您考虑可能影响吞吐量的内存、CPU 和其他系统资源，而前者只需您为表指定目标使用率、最小和最大容量。DynamoDB 可预置资源以实现读取和写入容量的目标使用率，然后根据使用情况自动扩展您的容量。如果您希望手动管理表吞吐量，则可以选择直接指定读取和写入容量。

下表总结了 DynamoDB 的关键定价概念。

资源类型	详细信息	每月价格
预置吞吐量 (写入)	一个写入容量单位 (WCU) 每秒最多执行一次写入操作，每月足以处理 250 万次写入操作	每 WCU 的价格低至 0.47 USD
预置吞吐量 (读取)	一个读取容量单位 (RCU) 每秒最多执行两次读取操作，每月足以处理 520 万次读取操作	每 RCU 的价格低至 0.09 USD
带索引的数据存储	DynamoDB 按每 GB 小时费率向您的表使用的磁盘空间收费	每 GB 的价格低至 0.25 USD

使用：
<https://us-east-2.console.aws.amazon.com/dynamodb/home?region=us-east-2#create-table>

Create DynamoDB table

Tutorial ?

DynamoDB is a schema-less database that only requires a table name and primary key. The table's primary key is made up of one or two attributes that uniquely identify items, partition the data, and sort data within each partition.

Table name*

i

Primary key*

Partition key

String

⌵

i

☐ Add sort key

Table settings

Default settings provide the fastest way to get started with your table. You can modify these default settings now or after your table has been created.

☒ Use default settings

No secondary indexes.

Provisioned capacity set to 5 reads and 5 writes.

Basic alarms with 80% upper threshold using SNS topic "dynamodb".

On-Demand Backup and Restore Enabled

NEW!

建表的时候可以自定义需求，后期也能修改，或者有个auto scaling

table has been created.

☐ Use default settings

Secondary indexes

Name	Type	Partition key	Sort key
+ Add index			

Provisioned capacity

Read capacity units
Table 5
Estimated cost \$2.91 / month ([Capacity calculator](#))

Capacity calculator

Avg. item size	1	KB
Item read/sec	1	Eventually consistent
Item write/sec	1	
Read capacity	1	
Write capacity	1	
Estimated cost per table/index	\$0.59 / month	
<button>Update</button>		

Auto Scaling

开发

底层接口提供多种语言：

python底层接口文档


<http://boto3.readthedocs.io/en/latest/guide/dynamodb.html>

https://docs.aws.amazon.com/zh_cn/amazondynamodb/latest/developerguide/GettingStarted.Python.html

dynamodb高层接口只看到提供了java sdk (<https://docs.aws.amazon.com/amazondynamodb/latest/developerguide/DynamoDBMapper.html>)

[AWS Documentation](#) » [Amazon DynamoDB](#) » [Developer Guide](#) » [Programming with DynamoDB and the AWS SDKs](#) » [Higher-Level Programming Interfaces for DynamoDB](#) » [Java: DynamoDBMapper](#)

Java: DynamoDBMapper



HBase

HBase is an open-source, non-relational, distributed database modeled after Google's Bigtable and is written in **Java**. It is developed as part of Apache Software Foundation's

Apache **Hadoop** project and runs on top of HDFS (**Hadoop** Distributed File System), providing Bigtable-like capabilities for **Hadoop**.

Storage Mechanism in HBase

HBase is a **column-oriented database** and the tables in it are sorted by row. The table schema defines only column families, which are the key value pairs. A table have multiple column families and each column family can have any number of columns. Subsequent column values are stored contiguously on the disk. Each cell value of the table has a timestamp. In short, in an HBase:

- Table is a collection of rows.
- Row is a collection of column families.
- Column family is a collection of columns.
- Column is a collection of key value pairs.

Given below is an example schema of table in HBase.

[illegible]

Features of HBase

- HBase is linearly scalable.
- It has automatic failure support.
- It provides consistent read and writes.
- It integrates with Hadoop, both as a source and a destination.
- It has easy java API for client.
- It provides data replication across clusters.

Where to Use HBase

- Apache HBase is used to have random, real-time read/write access to Big Data.
- It hosts very large tables on top of clusters of commodity hardware.
- Apache HBase is a non-relational database modeled after Google's Bigtable. Bigtable acts up on Google File System, likewise Apache HBase works on top of Hadoop and HDFS.

方案：

1. 数据规模暂时不大的话，可以试试dynamodb的免费套餐，或者拿小部分数据试试这个？(昨天建aws账号的时候，瞎点了两下，被划走二十多刀的我，已经有点阴影了，哈哈哈)

免费套餐

开始免费使用 DynamoDB。许多应用程序都可以在这些免费套餐限制内运行。

- 吞吐量每月足以处理多达 2 亿个请求 (25 个写入容量单位和 25 个读取容量单位)。
- 25GB 带索引的数据存储。
- 每月 250 万个 [DynamoDB Streams](#) 读取请求。
- 可在最多 2 个 AWS 地区部署 [DynamoDB Global Tables](#)。

您仅需为超出这些免费套餐限制的资源付费。DynamoDB 的免费套餐在 [AWS 免费套餐](#) 的 12 个月期限到期后不会过期，并且可供所有 AWS 客户使用。

2. 搭HBase集群，看下效果，稳定性，读写延迟什么的，人力成本估计会比较大，这两天（周六前）我试搭下HBase集群，如果能用的话，先用着？周日再找部分数据试用下dynamodb? btw, 公司里有这方面经验的人么？